

Higgsfield AI Prompt Skill

User Guide

A plain-English guide for getting the most out of this skill.

No coding knowledge required.

Teaches Claude how to write professional-quality prompts for
Higgsfield AI -- the cinematic video and image generation platform.

v3.7.0 | 2026-05-04 | Built by O-Side Media

Table of Contents

- 1. What Is This?
- 2. What Can It Do?
- 3. How to Install
- 4. Quick Start -- Your First Prompt
- 5. The MCSLA Formula
- 6. Choosing a Model
- 7. Generation Types
- 8. Cinema Studio 2.5
- 9. Cinema Studio 3.0 (Business/Team Plan)
- 10. Prompting Best Practices (Seedance 2.0)
- 11. Soul ID -- Character Consistency
- 12. Character Sheet Creation
- 13. Identity vs. Motion Separation
- 14. Genre Recipes
- 15. Genre Templates
- 16. Cinematic Image Prompts
- 17. Negative Constraints Reference
- 18. Troubleshooting
- 19. Top Tips
- 20. Memory System (Advanced)
- 21. Cinema Studio Advanced Features
- 22. Repository Contents
- 23. FAQ

v3.0

v3.0

1. What Is This?

This is a Claude skill -- a set of instructions that teaches Claude (Anthropic's AI) how to write professional-quality prompts for Higgsfield AI, a cinematic video and image generation platform.

Instead of learning Higgsfield's dozens of models, camera controls, motion presets, and prompt tricks yourself, you just tell Claude what you want in normal language and it writes an optimized, ready-to-paste prompt for you.

Think of it as hiring a cinematographer who knows every Higgsfield feature by heart.

2. What Can It Do?

- Write production-ready prompts for any Higgsfield video or image model
- Recommend the best model for your specific shot
- Guide you through Cinema Studio 2.5's professional multi-shot workflow
- Guide you through Cinema Studio 3.0's Smart mode, native audio, and @ references (Business/Team plan)
- Apply Seedance 2.0 prompting best practices: Intent over Precision, Genre Router, I2V Gate, Anti-Slop check
- Use 3D Mode (Gaussian Splatting) to explore scenes from new angles
- Generate batch variations with Grid Generation (2x2 and 4x4)
- Help you maintain character consistency with Soul ID and Character Sheets
- Apply named camera controls and motion presets
- Troubleshoot failed or poor-quality generations with a diagnostic tree
- Optimize your credit usage
- Build multi-shot sequences with per-character emotions
- Use Frame Extraction Loop for iterative creative workflows
- Insert objects and people not in your original start frame
- Automatically prevent common AI artifacts using shared negative constraints
- Separate identity from motion to prevent character drift during camera moves
- Match your request to genre templates with Cinema Studio 3.0 genre mappings and prompt length targets
- Describe audio as a first-class element with SCELA integration (BGM, SFX, dialogue)

3. How to Install

Option A -- Claude Cowork (Desktop App)

Drop the entire repo folder into your Cowork workspace. The skill loads automatically.

Option B -- Claude Code (Terminal)

```
git clone https://github.com/OSideMedia/higgsfield-ai-prompt-skill
cp -r higgsfield-ai-prompt-skill ~/.claude/skills/higgsfield
```

Option C -- Claude.ai (Web / Projects)

- Create a new Project in Claude.ai
- Upload the main SKILL.md file as your project instruction base
- Upload individual sub-skill files from the skills/ directory as project documents

4. Quick Start -- Your First Prompt

Just tell Claude what you want. Be as casual or specific as you like.

Casual request:

```
"Write me a Higgsfield prompt for a woman walking through a foggy forest at dawn"
```

Claude will respond with a complete, ready-to-paste prompt including the model recommendation, camera movement, style, and all the details.

More specific request:

```
"I need a Kling 3.0 prompt for a close-up dialogue scene between two characters
in a coffee shop. Warm lighting, shallow depth of field, handheld camera."
```

5. The MCSLA Formula

Every prompt is built on five layers -- the cinematographer's checklist:

Letter	Layer	Example
M	Model	Kling 3.0
C	Camera	FPV Drone weaving through the alley
S	Subject	A woman in a tactical jacket
L	Look	Cinematic, cold blue shadows, 16:9
A	Action	She sprints, slides under a gate

You don't need to specify all five layers every time. Claude fills in sensible defaults for anything you leave out.

6. Choosing a Model

Video Models

What you're making	Best model
Character-driven drama, dialogue	Kling 3.0
Epic scale, big environments	Sora 2
Lip-sync, multilingual dialogue	Seedance 1.5 Pro
Complex choreography, reference-based	Seedance 2.0
Long takes, camera control	Kling Motion Control
Motion transfer from reference video	Kling 3.0 Motion Control
60fps, first+last frame, reference images	Wan 2.7
Budget Veo 3.1 quality at volume	Veo 3.1 Lite
Quick iteration, drafts	Kling 2.5 Turbo
Nature, documentary	Veo 3 / Veo 3.1
Dance, fluid body motion	Minimax Hailuo 02
Cinema Studio 3.0 workflow	Business/Team plan

Image Models

What you're making	Best model
Portraits, fashion	Soul 2.0
Maximum photorealism	Nano Banana Pro
Fast pro-quality + text	Nano Banana 2
Text/logo rendering	GPT Image 1.5
Reference editing	Seedream 4.5
Cinematic keyframes for I2V	Soul Cinema Preview

7. Generation Types

Text-to-Video (T2V)

Describe the scene from scratch. No input image needed. Best for establishing shots, environments, abstract concepts.

Image-to-Video (I2V)

Upload a still image and describe what should move or change. Best for character consistency, product shots, bringing storyboards to life.

Key rule for I2V: Do NOT re-describe what's already in the image. Only describe what should change or animate. (See I2V Gate rule in Section 10.)

8. Cinema Studio 2.5

Cinema Studio is Higgsfield's professional filmmaking environment. It gives you control over optical physics, multi-shot sequences, character elements, Soul Cast AI actors, and built-in color grading.

The 10-step workflow:

Step	What you do	Details
1	Script	Write your scene description / shot list
2	Soul Cast	Generate AI actors from parameters (no photos needed)
3	Reference	Upload character photo or use Soul Cast actor as Reference Anchor
4	Elements	Define @Characters, @Locations, @Props
5	Optical Stack	Choose camera body + lens + focal length + aperture
6	Hero Frame	Generate a key image to lock the visual tone
7	Color Grade	Apply color grading to keyframes before video generation
8	Camera Config	Set Director Panel movement + Speed Ramp + Duration
9	Shot Mode	Single Shot, Multi-Shot Auto, or Multi-Shot Manual
10	Generate	Run generation and export

Director Panel: 18 camera movements -- Static, Handheld, Zoom Out/In, Camera Follows, Pan Left/Right, Tilt Up/Down, Orbit Around, Dolly In/Out, Jib Up/Down, Drone Shot, Dolly Left/Right, 360 Roll, Auto.

Speed Ramp: Linear, Slow Mo, Speed Up, Impact, Auto, Custom (with editable curve).

Shot structure: Up to 6 scenes, 12s total max, per-scene config.

9. Cinema Studio 3.0 (Business/Team Plan)

v3.0

Cinema Studio 3.0 is available exclusively on Business and Team plans. Free and individual plan users should use Cinema Studio 2.5, which remains fully supported.

Cinema Studio 3.0 is a separate generation engine from 2.5. Switch between them using the version selector in the upper-right corner of the Cinema Studio UI. Both versions coexist -- Cinema Studio 2.5 is NOT deprecated.

Resolution Comparison Table

Feature	Cinema Studio 2.5	CS 3.0 (Biz/Team)
Video Resolution	Up to 1080p	Up to 720p (may increase)
Image Resolution	Up to 4K	Up to 2K
Max Duration	12s	15s
Aspect Ratios	6 options	7 (+ 21:9 ultrawide)
Audio	On/Off	On/Off (native stereo)
Shot Control	Manual multi-shot	Smart + Custom
Generation Cost	Varies	48 credits

Resolution limits for Cinema Studio 3.0 are subject to change. If you need higher resolution now, use Cinema Studio 2.5.

Key Differences in 3.0

- Native dual-channel stereo audio -- generated simultaneously with video, not post-processed
- Smart shot control -- model auto-plans camera based on genre and scene description
- 21:9 ultrawide aspect ratio (not available in 2.5)
- Up to 15s video duration (vs 2.5's 12s)
- 7 genres: General, Action, Horror, Comedy, Noir, Drama, Epic
- 7 Speed Ramp presets: Auto, Slow-mo, Ramp Up, Flash In, Flash Out, Bullet Time, Hero Moment
- No optical physics engine, color grading suite, 3D Mode, or Grid Generation (use 2.5 for these)

@ Reference Input Limits

Type	Max Count	Formats	Limit
Images	9	jpeg/png/webp/bmp	--
Video clips	3	mp4/mov	Combined <=15s
Audio clips	3	mp3/wav	Combined <=15s
Total files	<=12		

When to Use 3.0 vs 2.5

Need	Recommendation
Highest resolution (1080p/4K)	Cinema Studio 2.5

Longer duration (up to 15s)	Cinema Studio 3.0
Native audio with video	Cinema Studio 3.0
Optical physics / color grading	Cinema Studio 2.5
3D Mode / Grid Generation	Cinema Studio 2.5
Ultrawide 21:9	Cinema Studio 3.0
Smart auto-camera	Cinema Studio 3.0
Free/Individual plan	Cinema Studio 2.5

Cinema Studio 3.5 -- Briefly

Cinema Studio 3.5 sits alongside 2.5 and 3.0 in the model selector -- all three coexist, version is user-selected, no auto-routing. 3.5 collapses creative control into a three-pill main UI (Genre / Style / Camera), each defaulting to Auto with manual override. Camera Settings restores optical physics via a four-axis panel (Camera Body / Lens / Focal Length / Aperture, with 75mm added). Style Settings exposes three preset axes (Color Palette / Lighting / Camera Moveset Style) plus a free-form Manual Style mode. Image mode adds a Cinematic models picker with four options: Soul Cinema (default), Cinematic Characters, Cinematic Locations, and Cinematic Cameras (the latter surfaces the 2.5 optical vocabulary).

For full Cinema Studio 3.5 documentation -- three-pill UI, Style Settings, Camera Settings, Manual Style mode, Image Mode, four Cinematic models picker, Physics Rendering Decision Matrix -- see the higgsfield-cinema sub-skill.

10. Prompting Best Practices (Seedance 2.0)

v3.0

These best practices apply to Cinema Studio 3.0's generation engine and complement the MCSLA formula. They are not a replacement -- use MCSLA as the primary framework, then apply these refinements.

Intent over Precision

Tell the model WHAT you want and HOW it should FEEL, not every micro-detail. Short prompts (30-100 words) consistently outperform long ones. The model is an AI director you collaborate with, not a render engine you command.

Genre Router -- Prompt Length Targets

Genre	Lead With	Target Length
Product/E-commerce	Subject	30-50 words
Lifestyle/Social	Action	40-60 words
Drama/Narrative	Scene	60-100 words
Music Video	Style	50-80 words
Landscape/Travel	Scene	30-60 words
Commercial/Brand	Style	40-70 words
Anime/Artistic	Style	50-90 words

I2V Gate Rule

When using image references (@Image), describe ONLY motion and camera movement. NEVER re-describe what's already visible in the image. The model can see the image -- re-describing creates conflict and degrades output.

Anti-Slop Vocabulary

Kill these words -- they add zero information and waste tokens:

Kill	Replace With
beautiful/stunning	(delete -- describe specific visual)
epic	large-scale, sweeping, towering
amazing	(delete -- show, don't tell)
dynamic	fast-tracking, whip-pan, handheld
cinematic camera movement	slow dolly push / crane up / tracking

The One-Move Rule

For any single shot, specify only ONE primary camera move. Do NOT stack multiple moves (e.g., dolly push + pan left + tilt up). This is the #1 cause of jitter, unwanted rotation, and failed generations.

Physics Language

Use concrete physics consequences instead of mood words. Not 'powerful punch' but 'fist connects, sweat flies off in slow motion, opponent's head snaps back.' Not 'dramatic entrance' but 'door slams open, dust erupts from the frame,

Audio as First-Class Element (SCELA)

Describe audio separately in prompts. BGM, ambient SFX, and dialogue are parallel tracks via dual-channel stereo generation. Specific foley descriptions directly influence output: 'the scratch of frosted glass, rustling plush fabric, gentle tapping on acrylic.'

No Negative Prompts

Cinema Studio 3.0 does not support negative prompt syntax. Use positive constraints: 'locked-off static camera' instead of 'no shaky camera.' 'sharp focus throughout' instead of 'no blur.'

11. Soul ID -- Character Consistency

Soul ID keeps a character looking the same across multiple generations. Upload a clear reference photo, create a Soul ID, and every future prompt can reference that same character.

Best Practices for the Reference Photo:

- Front-facing or 3/4 angle -- full face visible
- Even lighting -- no harsh shadows
- Neutral expression (slight smile is fine)
- Clear image -- no blur, no obstruction
- Solo subject -- no other people

Cinema Studio 3.0 Character Consistency

v3.0

- Use 2-3 clear, well-lit reference shots (frontal, 3/4-angle, side profile)
- Outfit descriptions must be specific -- materials, colors, distinctive details
- In I2V workflows: describe what the character DOES, not what they LOOK LIKE
- If features drift: use character sheet directly as @Image1 for tighter anchoring
- Multi-character scenes: reference each character separately with distinct @Image tags

12. Character Sheet Creation

A character sheet is a multi-angle reference image showing the same character from several viewpoints -- front, 3/4, side profile, and back. It gives Soul ID far more geometry data than a single photo.

- Generate your character using your preferred model
- Use Grid Generation (2x2 or 4x4) to produce multiple variations
- Use 3D Mode to orbit and capture front, side, and 3/4 angles
- Arrange the best angles into a single composite reference image
- Upload as your Soul ID reference

13. Identity vs. Motion Separation

When Soul ID or character consistency is involved, every prompt must be split into two clearly labeled blocks:

Identity Block -- Static visual descriptors ONLY

- Face features, skin tone, body type, distinguishing marks
- Clothing, accessories, color palette
- NO motion, NO camera, NO temporal language

Motion Block -- Temporal and camera ONLY

- Camera movement, action choreography, speed
- Environmental motion, atmospheric changes
- NO character appearance repetition

Example (Good -- separated):

```
Identity Block:
The Soul ID character -- sharp cheekbones, auburn hair shoulder-length,
wearing a blue trench coat with silver buttons, lean athletic build.

Motion Block:
She runs through a rain-soaked alley, coat flapping behind her.
Camera: Action Run -- low behind, matching pace.
Neon reflections streak across wet concrete.
Style: Cinematic, cold blue shadows, warm neon accents. 16:9.
```

14. Genre Recipes

Genre	Story Pattern
Action / Chase	Establish -> Pursuit -> Obstacle -> Climax
Emotional Drama	Context -> Tension -> Breaking Point -> Resolution
Horror	Calm -> Unease -> Build -> Scare
Product Ad	Hero Shot -> Feature Detail -> Lifestyle -> CTA
Sci-Fi	World -> Discovery -> Conflict -> Revelation
Romance	Meeting -> Tension -> Connection Moment
Documentary	Environment -> Subject in Context -> Observational
Dance / Music	Establish Space -> Performance Builds -> Beat Sync
Transformation	Before State -> Trigger -> Transform -> After State

15. Genre Templates

10 deeply annotated prompt templates in the templates/ folder. Each includes: when to use, recommended model, full example prompt, line-by-line annotation, negative constraints, common mistakes, variations, Identity/Motion blocks, and Cinema Studio 3.0 genre mappings with prompt length targets.

#	Template	CS 3.0 Genre	Prompt Length
01	Cinematic Action Chase	Action	60-100 words
02	Product / UGC Showcase	General	30-50 words
03	Horror / Atmospheric Dread	Horror	60-100 words
04	Fashion / Editorial	Drama/General	40-60 words
05	Sci-Fi / VFX Spectacle	Epic/Action	50-90 words
06	Portrait / Character Intro	Drama	60-100 words
07	Landscape / Establishing	Epic/General	30-60 words
08	Comedy / Social Media	Comedy	40-60 words
09	Romantic / Intimate	Drama	60-100 words
10	Dance / Music Performance	Action/Drama	50-80 words

16. Cinematic Image Prompts

The skill includes a complete cinematic shot reference for still images -- 10 distance/size shots, 10 camera angles, and 17 camera movement keywords.

Image Prompt Formula:

[Shot size] + [Angle] + [Movement keyword] of [character].
[Pose]. [Environment]. [Lighting]. [Style].

17. Negative Constraints Reference

A shared constraints file consolidates all known AI generation artifacts and the prompt phrasing to prevent them.

Category	Covers	Check when
Body / Motion	Floating limbs, jittery motion	Action prompts
Face / Identity	Face morphing, identity drift	Character prompts
Texture / Lighting	Flickering, style ignored	Style-heavy prompts
Temporal	Static I2V, camera fails, lip-sync	Multi-shot, I2V
Content Filter	Blocked content, brands, persons	Horror, branded
Cinema Studio	512 char limit, @ Element bugs	CS 2.5 prompts
CS 3.0 Notes	No negative prompts -- positive only	CS 3.0 prompts

Cinema Studio 3.0 does not support negative prompt syntax. All constraints must be phrased as positive statements. The shared constraints file includes a positive alternatives table.

18. Troubleshooting

Problem	Quick Fix
Character face keeps changing	Soul ID + character sheet + Identity/Motion separation
Video is static / not animating	Describe ONLY what changes (I2V Gate rule)
Camera movement ignored	Use exact preset name (e.g. 'Dolly In')
Style not applying	Put style at end; use One Style Anchor Rule
Generation filtered	Check memory DB for known workarounds
Too many actions	Split into separate shots (One Action Per Shot)
Identity drift during moves	Separate Identity Block from Motion Block

Cinema Studio 3.0 Diagnostic Tree

v3.0		
Symptom	Likely Cause	Fix
Blurry/jittery/morphing	Overspecification	Cut to 30-100 words, use @ref
Camera chaotic	Multiple moves	ONE move per shot (One-Move Rule)
Character mismatch	Re-describing character	Delete appearance, keep action
Action stiff	No physics language	Add adverbs + consequences
Just not right	Ambiguous prompt	Run Anti-Slop Check
Audio mismatch	Conflicting audio desc	Timestamp anchor, remove SFX

19. Top Tips

1.

Be specific. Name camera presets, describe VFX concretely. 'Dolly In' beats 'the camera moves forward.'

2.

One action per shot. AI renders clean physics for one action. Chain multiple shots for complex sequences.

3.

Subject first, style last. Subject -> Action -> Camera -> Style is the most reliable prompt order.

4.

Keep it under 200 words. Focused prompts outperform exhaustive ones. For Cinema Studio 3.0: 30-100 words is the sweet spot.

5.

Cinema Studio 2.5: 512 character limit. @ Element chips consume ~80-100 hidden characters each.

6.

Use the Hero Frame method. Generate the perfect still image first, then animate it.

7.

Slow motion trick. If fast action keeps breaking, generate in Slow Mo and speed up in post.

8.

Don't say 'cinematic masterpiece.' Replace with actual visual details (One Style Anchor Rule).

9.

12-second cap. Multi-Shot Manual sequences max 12 seconds total across all scenes.

10.

Separate identity from motion. Keep face descriptors and camera/action in separate blocks.

11.

Check templates first. Before writing from scratch, see if a genre template matches your request.

12.

One-Move Rule. Specify only ONE primary camera move per shot to avoid jitter.

13.

Intent over Precision. Describe what you want and how it should feel, not every micro-detail.

14.

Audio is first-class. Describe BGM, SFX, and dialogue separately for Cinema Studio 3.0.

15.

Use @Video for complex motion. Camera transfer and choreography cloning via video reference is the most reliable method.

20. Memory System (Advanced)

The skill includes a memory system that stores what works and what doesn't. Content filter workarounds and quality failure fixes are stored in JSON databases and consulted automatically before generation.

You don't need to manage this yourself -- Claude checks the memory automatically when the recall skill is loaded.

21. Cinema Studio Advanced Features

Soul Cast -- AI Actor Generation (2.5 + 3.0)

Generate AI actors from 8 parameter categories (Genre, Budget, Era, Archetype, Identity, Physical Appearance, Details, Outfit). In Cinema Studio 3.0 (Business/Team): General (2K) / Character (4K) / Location (4K) modes, 0.125 credits per image.

Built-in Color Grading (2.5 only)

Color temperature, contrast, saturation, sharpness, film grain, exposure, bloom -- applied to keyframes before video generation. Not available in Cinema Studio 3.0.

3D Mode -- Gaussian Splatting (2.5 only)

Generate an image, then enter 3D Mode to build a 3D version. Orbit the virtual camera to explore from any angle. Not available in Cinema Studio 3.0.

Grid Generation -- Batch Variations (2.5 only)

Generate 2x2 (4 variations) or 4x4 (16 variations) from a single prompt. Not available in Cinema Studio 3.0.

Resolution Settings (2.5)

Explicit resolution control: 1K (fast drafts), 2K (default), 4K (final delivery). Cinema Studio 3.0: up to 720p video, up to 2K image.

Frame Extraction Loop (2.5)

Build an image -> Animate -> Extract a frame -> Feed it back -> Repeat.

Object and Person Insertion (2.5)

Insert characters or objects not in your original start frame.

Smart Shot Control (3.0 only)

Model auto-plans camera language based on genre and scene description. Trust it for genre-appropriate camera work. Override by switching to Custom multi-shot.

Native Audio-Video Joint Generation (3.0 only)

Audio generated simultaneously with video via unified multimodal architecture. Dual-channel stereo. Sound design prompts directly influence both audio AND visual generation.

22. Repository Contents

Root Files

File	What it is
SKILL.md	Model selection guide (routes model questions)
README.md	Installation and usage guide
CHANGELOG.md	Version history
USER-GUIDE.pdf	This document
model-guide.md	Model comparison tables + decision flowchart
image-models.md	Image model reference + pricing tiers
vocab.md	Full platform vocabulary reference
prompt-examples.md	High-quality example prompts

Sub-Skills (20 total)

Sub-Skill	What it covers
higgsfield-prompt	MCSLA formula + Seedance 2.0 best practices
higgsfield-image-shots	Cinematic image prompting
higgsfield-models	Model selection guide + CS 3.0 comparison
higgsfield-camera	Camera controls + One-Move Rule + Smart Mode
higgsfield-motion	Motion presets + intent-first choreography
higgsfield-style	Visual styles + One Style Anchor Rule
higgsfield-soul	Soul ID + Soul Cast 3.0 + Soul Cinema (CS 3.0/3.5 default image model)
higgsfield-audio	Audio prompting + CS 3.0 native audio (SCELA)
higgsfield-apps	One-click Apps guide (80+)
higgsfield-recipes	Genre templates
higgsfield-troubleshoot	Fix generations + CS 3.0 diagnostic tree
higgsfield-assist	Platform copilot + credit optimization
higgsfield-mixed-media	Artistic preset overlays
higgsfield-moodboard	Moodboard + style consistency
higgsfield-pipeline	Multi-step production pipeline
higgsfield-cinema	Cinema Studio 2.5 + 3.0 + 3.5 (Soul Cast, Image Mode, Cinematic models)
higgsfield-recall	Pre-generation memory check
higgsfield-vibe-motion	Vibe Motion / motion graphics
higgsfield-seedance	Seedance 2.0 prompt director + content-filter preflight
higgsfield-workspaces	Workspace-first decision layer

23. FAQ

Q: Do I need a Higgsfield account?

A: Yes -- this skill writes prompts for you, but you paste and run them on higgsfield.ai.

Q: Which Claude plan do I need?

A: Any plan works -- Free, Pro, or Team. The skill loads as project instructions.

Q: Do I need a Business/Team plan for Cinema Studio 3.0?

A: Yes -- Cinema Studio 3.0 is exclusive to Business and Team plans on Higgsfield. Cinema Studio 2.5 is available on all plans.

Q: Can I use this with other AI video tools?

A: The prompts are optimized for Higgsfield specifically. General prompt techniques (MCSLA, shot framing) transfer to other tools.

Q: How do I get updates?

A: The skill is versioned (currently v3.7.0). Check the repository for updates.

Q: Can I contribute?

A: Yes! Fork the repo, add your improvements, and submit a pull request.

Q: What changed since v3.0.0?

A: Seven platform releases shipped between v3.0.0 (April 2026) and this guide (v3.6.2). The major themes: install-path simplification (v3.3.0), workspace-first decision layer + Seedance 2.0 prompt modes + Kling 3.0 Motion Control + Location Reference Sheets (v3.4.0), Reference Sheet Types + Iteration Rule (v3.4.1), workspace expansion (v3.5.0), Cinema Studio 3.5 with Image Mode + Elements System library surface + Physics Rendering Decision Matrix (v3.6.0), markdownlint config (v3.6.1), and this documentation refresh (v3.6.2). See CHANGELOG.md and the higgsfield-cinema sub-skill for full per-release detail.

Built by O-Side Media | v3.7.0 | 2026-05-04 | Platform: higgsfield.ai